



## How To Set Up MySQL Master-Master Replication

Posted April 26, 2013

 739.8k

MYSQL

SCALING

BACKUPS



60

By: Jason Kurtz

### Status: Deprecated

This tutorial has been deprecated and is no longer maintained.

**Reason:** This tutorial covers an outdated method for building a multi-master topology. At the time this article was originally published, group replication had not yet been implemented in MySQL.

#### See Instead:

You can read the newer [How To Configure MySQL Group Replication on Ubuntu 16.04](#) tutorial to set up a multi-primary replication group.

### Intro

This second installment of "Scaling Web Applications" will list out the steps necessary for scaling a mysql deployment over two VPS.

The first article in this series laid out the steps needed to load-balance nginx over two VPS, and it is recommended that you read [that article](#) first.

MySQL replication is the process by which a single data set, stored in a MySQL database, will be live-copied to a second server. This configuration, called "master-slave" replication, is a typical setup. Our setup will be better than that, because master-master replication allows data to be copied from either server to the other one. This subtle but important difference allows us to perform mysql read or writes from either server. This configuration adds redundancy and increases efficiency when dealing with accessing the data.

The examples in this article will be based on two VPS, named Server C and Server D.

### Server C: 3.3.3.3

### Server D: 4.4.4.4

## Step 1 - Install and Configure MySQL on Server C

The first thing we need to do is to install the mysql-server and mysql-client packages on our server. We can do that by typing the following:

```
sudo apt-get install mysql-server mysql-client
```

By default, the mysql process will only accept connections on localhost (127.0.0.1). To change this default behavior and change a few other settings necessary for replication to work properly, we need to edit /etc/mysql/my.cnf on Server C. There are four lines that we need to change, which are currently set to the following:

```
#server-id          = 1
#log_bin            = /var/log/mysql/mysql-bin.log
#binlog_do_db       = include_database_name
bind-address        = 127.0.0.1
```

The first of those lines is to uniquely identify our particular server, in our replication configuration. We need to uncomment that line, by removing the "#" before it. The second line indicates the file in which changes to any mysql database or table will be logged.

The third line indicates which databases we want to replicate between our servers. You can add as many databases to this line as you'd like. The article will use a single database named "example" for the purposes of simplicity. And the last line tells our server to accept connections from the internet (by not listening on 127.0.0.1).

```
server-id          = 1
log_bin            = /var/log/mysql/mysql-bin.log
binlog_do_db       = example
# bind-address      = 127.0.0.1
```

Now we need to restart mysql:

```
sudo service mysql restart
```

We next need to change some command-line settings within our mysql instance. Back at our shell, we can get to our root mysql user by typing the following:

```
mysql -u root -p
```

Please note that the password this command will prompt you for is that of the root mysql user, not the root user on our droplet. To confirm that you are logged in to the mysql shell, the prompt should look like the following.

```
mysql>
```

Once we are logged in, we need to run a few commands.

We need to create a pseudo-user that will be used for replicating data between our two VPS. The examples in this article will assume that you name this user "replicator". Replace "password" with the password you wish to use for replication.

```
create user 'replicator'@'%' identified by 'password';
```

Next, we need to give this user permissions to replicate our mysql data:

```
grant replication slave on *.* to 'replicator'@'%';
```

Permissions for replication cannot, unfortunately, be given on a per-database basis. Our user will only replicate the database(s) that we instruct it to in our config file.

For the final step of the initial Server C configuration, we need to get some information about the current MySQL instance which we will later provide to Server D.

The following command will output a few pieces of important information, which we will need to make note of:

```
show master status;
```

The output will look similar to the following, and will have two pieces of critical information:

```
+-----+-----+-----+-----+
| File           | Position | Binlog_Do_DB | Binlog_Ignore_DB |
+-----+-----+-----+-----+
| mysql-bin.000001 |      107 | example      |                   |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

We need to make a note of the file and position which will be used in the next step.

## Step 2 - Install and Configure MySQL on Server D

We need to repeat the same steps that we followed on Server C. First we need to install it, which we can do with the following command:

```
sudo apt-get install mysql-server mysql-client
```

Once the two packages are properly installed, we need to configure it in much the same way as we configured Server C. We will start by editing the /etc/mysql/my.cnf file.

```
sudo nano /etc/mysql/my.cnf
```

We need to change the same four lines in the configuration file as we changed earlier.

The defaults are listed below, followed by the changes we need to make.

```
#server-id          = 1
#log_bin            = /var/log/mysql/mysql-bin.log
#binlog_do_db       = include_database_name
bind-address        = 127.0.0.1
```

We need to change these four lines to match the lines below. Please note, that unlike Server C, the server-id for Server D cannot be set to 1.

```
server-id           = 2
log_bin             = /var/log/mysql/mysql-bin.log
binlog_do_db        = example
# bind-address       = 127.0.0.1
```

After you save and quit that file, you need to restart mysql:

```
sudo service mysql restart
```

It is time to go into the mysql shell and set some more configuration options.

```
mysql -u root -p
```

First, just as on Server C, we are going to create the pseudo-user which will be responsible for the replication. Replace "password" with the password you wish to use.

```
create user 'replicator'@'%' identified by 'password';
```

Next, we need to create the database that we are going to replicate across our VPS.

```
create database example;
```

And we need to give our newly created 'replication' user permissions to replicate it.

```
grant replication slave on *.* to 'replicator'@'%';
```

The next step involves taking the information that we took a note of earlier and applying it to our mysql instance. This will allow replication to begin. The following should be typed at the mysql shell:

```
slave stop;
CHANGE MASTER TO MASTER_HOST = '3.3.3.3', MASTER_USER = 'replicator', MASTER_PASSWORD = 'password',
slave start;
```

You need to replace 'password' with the password that you have chosen for replication. Your values for MASTERLOGFILE and MASTERLOGPOS may differ than those above. You should copy the values that "SHOW

MASTER STATUS" returns on Server C.

The last thing we have to do before we complete the mysql master-master replication is to make note of the master log file and position to use to replicate in the other direction (from Server D to Server C).

We can do that by typing the following:

```
SHOW MASTER STATUS;
```

The output will look similar to the following:

```
+-----+-----+-----+-----+
| File           | Position | Binlog_Do_DB | Binlog_Ignore_DB |
+-----+-----+-----+-----+
| mysql-bin.000004 |      107 | example      |                   |
+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Take note of the file and position, as we will have to enter those on server C, to complete the two-way replication.

The next step will explain how to do that.

## Step 3 - Completing Replication on Server C

Back on Server C, we need to finish configuring replication on the command line. Running this command will replicate all data from Server D.

```
slave stop;
CHANGE MASTER TO MASTER_HOST = '4.4.4.4', MASTER_USER = 'replicator', MASTER_PASSWORD = 'password',
slave start;
```

Keep in mind that your values may differ from those above. Please also replace the value of MASTER\_PASSWORD with the password you created when setting up the replication user.

The output will look similar to the following:

```
Query OK, 0 rows affected (0.01 sec)
```

The last thing to do is to test that replication is working on both VPS. The last step will explain an easy way to test this configuration.

## Step 4 - Testing Master-Master Replication

Now that have all the configuration set up, we are going to test it now. To do this, we are going to create a table in our example database on Server C and check on Server D to see if it shows up. Then, we are going to delete it from Server D and make sure it's no longer showing up on Server C.

We now need to create the database that will be replicated between the servers. We can do that by typing the following at the mysql shell:

```
create database example;
```

Once that's done, let's create a dummy table on Server C:

```
create table example.dummy (`id` varchar(10));
```

We now are going to check Server D to see if our table exists.

```
show tables in example;
```

We should see output similar to the following:

```
+-----+  
| Tables_in_example |  
+-----+  
| dummy             |  
+-----+  
1 row in set (0.00 sec)
```

The last test to do is to delete our dummy table from Server D. It should also be deleted from Server C.

We can do this by entering the following on Server D:

```
DROP TABLE dummy;
```

To confirm this, running the "show tables" command on Server C will show no tables:

```
Empty set (0.00 sec)
```

And there you have it! Working mysql master-master replication. As always, any feedback is more than welcome.

By: Jason Kurtz

♡ Upvote (60)

📌 Subscribe

🔗 Share

Build something great with a \$100, 60 day credit

Build the internet on DigitalOcean with a \$100, 60 day credit to use across Droplets, Block Storage, Load Balancers and more!

REDEEM CREDIT

## Related Tutorials

How To Set Up a Remote Database to Optimize Site Performance with MySQL on Ubuntu 18.04

How To Set Up Laravel, Nginx, and MySQL with Docker Compose

An Introduction to Queries in MySQL

How To Manage an SQL Database

How To Install Linux, Nginx, MySQL, PHP (LEMP stack) on Debian 9

## 143 Comments

Leave a comment...

Log In to Comment

^ [marlon](#) June 20, 2013



0 This is excellent, Followed it by example, then set up my master-master-slave from there. Thanks D.O.

^ [kamaln7](#) MOD June 20, 2013



0 @Marlon: Thanks for the awesome feedback! I'm really glad to hear that this article was helpful to you :]

^ [marlon](#) June 20, 2013



3 Karmal, cool!

I hate posting again, but I forgot something. I could be way off, but I think adding the `auto_increment` and `auto_increment_offset` directives might be a good idea for master-master, as this could save folks from potential PK conflicts, or users could investigate this for themselves: [http://dev.mysql.com/doc/refman/5.7/en/replication-options-master.html#sysvar\\_auto\\_increment\\_increment](http://dev.mysql.com/doc/refman/5.7/en/replication-options-master.html#sysvar_auto_increment_increment)

So on Server C: 2/1 and for server D: 2/2 or something like that.

^ [offgl](#) July 29, 2017



0 Correct this doc is useless and I find most of digital ocean documentations are not tested before posting. Just some useless fellow copying from another website and changing few words and posting it here.

^ [rmang](#) July 3, 2013



<sup>1</sup> How about if you have an existing database and want to set up master-master? What needs to be changed for this scenario?

---

^ kamaln7 MOD July 3, 2013

0 @rmang: This article should still work even if you have an already existing mysql setup.

---

^ dave119711 July 4, 2013

0 These guys have a really good replication configurator <http://www.severalnines.com/>

---

^ vguhesan+do July 29, 2013

0 what is the minimal MySQL version that supports this model? 5.0, 5.1, 5.5, 5.6???

---

^ kamaln7 MOD July 29, 2013

0 @vguhesan+do: Replication capabilities enabling the databases on one MySQL server to be duplicated on another were introduced in MySQL 3.23.15.

---

^ masud August 28, 2013

0 Kamal, great article. Congratulations.

Now I have one scenario on hand. We have 2 servers under one physical load balancer. Its configured using round robin method. Now we know the content / file replication or synchronization wont be an issue but can we have the same database on both dedicated servers ( under same network! ) and do real-time replication using your method or you suggest something else?

We did consider rsync for mysql as well but it can be tricky as the db has to be stopped while replicating and all.

Can you please advice?

Thanks in advance mate.

---

^ kamaln7 MOD August 28, 2013

0 @masud: I recommend using one separate database server and having it accessible by both web servers so that they access the exact same data.

---

^ sarayut\_utsakoo September 13, 2013

0 This is a really good article what I'm look for.  
Thanks for share.

---

^ m.shoaib.mir September 17, 2013

0 Is this article work for two droplets ? I have been trying to configure master master on two droplets with different ips and db servers. But it doesn't work. All is set and no data is replicated. Any help ?

---

kamaln7 MOD September 17, 2013



^ @m.shoaib.mir: Check mysql's binary logs:

<http://egloo.wordpress.com/2008/11/19/how-to-read-mysql-binary-logs/> - see if you can find any errors/pointers there.

---

^ m.shoaib.mir September 22, 2013

Hi Kamal,

I am getting the following error in Master Master Replication between two droplets. I have found the error through the command "mysql> show slave status\G;"

Error - error connecting to master 'replicator@Droplet IP:3306'

It means that Master-Master or Master Slave are not connecting with each other.

Help!!!

---

^ kamaln7 MOD September 23, 2013

@m.shoaib.mir: Did you follow steps 1 and 2?

---

^ m.shoaib.mir September 25, 2013

@Kamal Naseer

Yes, I had followed steps 1 and 2.

---

^ hi.suketu82 September 25, 2013

Which edition of MYSQL is required for this setup. Does it works in Standard edition?

---

^ kamaln7 MOD September 25, 2013

@hi.suketu82: It should work in all editions of MySQL.

---

^ kamaln7 MOD September 25, 2013

@m.shoaib.mir: Are the droplets firewalled? What's the output of

```
sudo iptables -L -n -v
```

on each droplet?

---

^ m.shoaib.mir September 26, 2013

//////////////////////////////// Droplet 1 //////////////////////////////////

Chain INPUT (policy ACCEPT 679 packets, 46247 bytes)  
pkts bytes target prot opt in out source destination

Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)  
pkts bytes target prot opt in out source destination

Chain OUTPUT (policy ACCEPT 732 packets, 59559 bytes)

[illegible]

Note: i was already doing all the steps

